

Diagnosing Agent Failures at Scale

Using an LLM judge to automatically label why AI agents fail

Huy Huynh, Aditya Kumar, Ava Doan · Advised by Tapan Chugh & Arvind Krishnamurthy



University of Washington

RESULT

0.5573 → 0.7405

Hierarchical judging closes most of the gap between a single-pass LLM judge and human annotation.

micro-F1 on BFCL hold-out · 35 GPT-5 traces · 142 label instances · macro-F1 also improves from 0.52 to 0.62

PROBLEM STATEMENT

Benchmark scores alone don't help us understand *why* agents are failing.

- An **agent** is an LLM (e.g., GPT-5) wired up to tools such as Gmail, GitHub, or a file system, so it can take actions on the user's behalf instead of just chatting.
- When the agent fails, the benchmark only reports **valid: False**. That tells us nothing about the cause. Was the wrong tool called? Was a step skipped? Did the agent invent a parameter? Each case needs a different fix.

REAL FAILURE · BFCL · GPT-5

user: "Write the statistics to Annual_Report_2023.docx: *Earning: 2000, Expenditure: 500, Name: Gorilla*"
agent: reorders fields, writes them on separate lines
↳ **valid: False · instance_state_mismatch**
The agent finished. Just not exactly as required. "Failed" is the wrong word; "completed but did not match" is the right one.

- The AFT source annotations hand-read 500+ such failures and built the **Agent Failure Taxonomy (AFT)**, a hierarchical vocabulary of 19 failure categories. Hand-annotation does not scale. **We built an LLM judge** that reproduces those labels automatically across the 15 categories we focus on (we exclude benchmark-bug labels for now).

TAXONOMY DEFINITIONS


AFT categorizes all agent failures into 19 distinct failure categories.

AFT organizes failures along two dimensions: **completion status** (did the agent finish?) and **workflow phase** (intent, reasoning, or output). Our *Annual Report* case sits in the bottom-right cell, an **Output State Mismatch**.

EXAMPLE FAILURE TYPE

Missing Tacit Knowledge

The agent lacked basic knowledge the task assumed but did not state — e.g., asked to "cancel my booking," it never looked up the user's booking ID first.

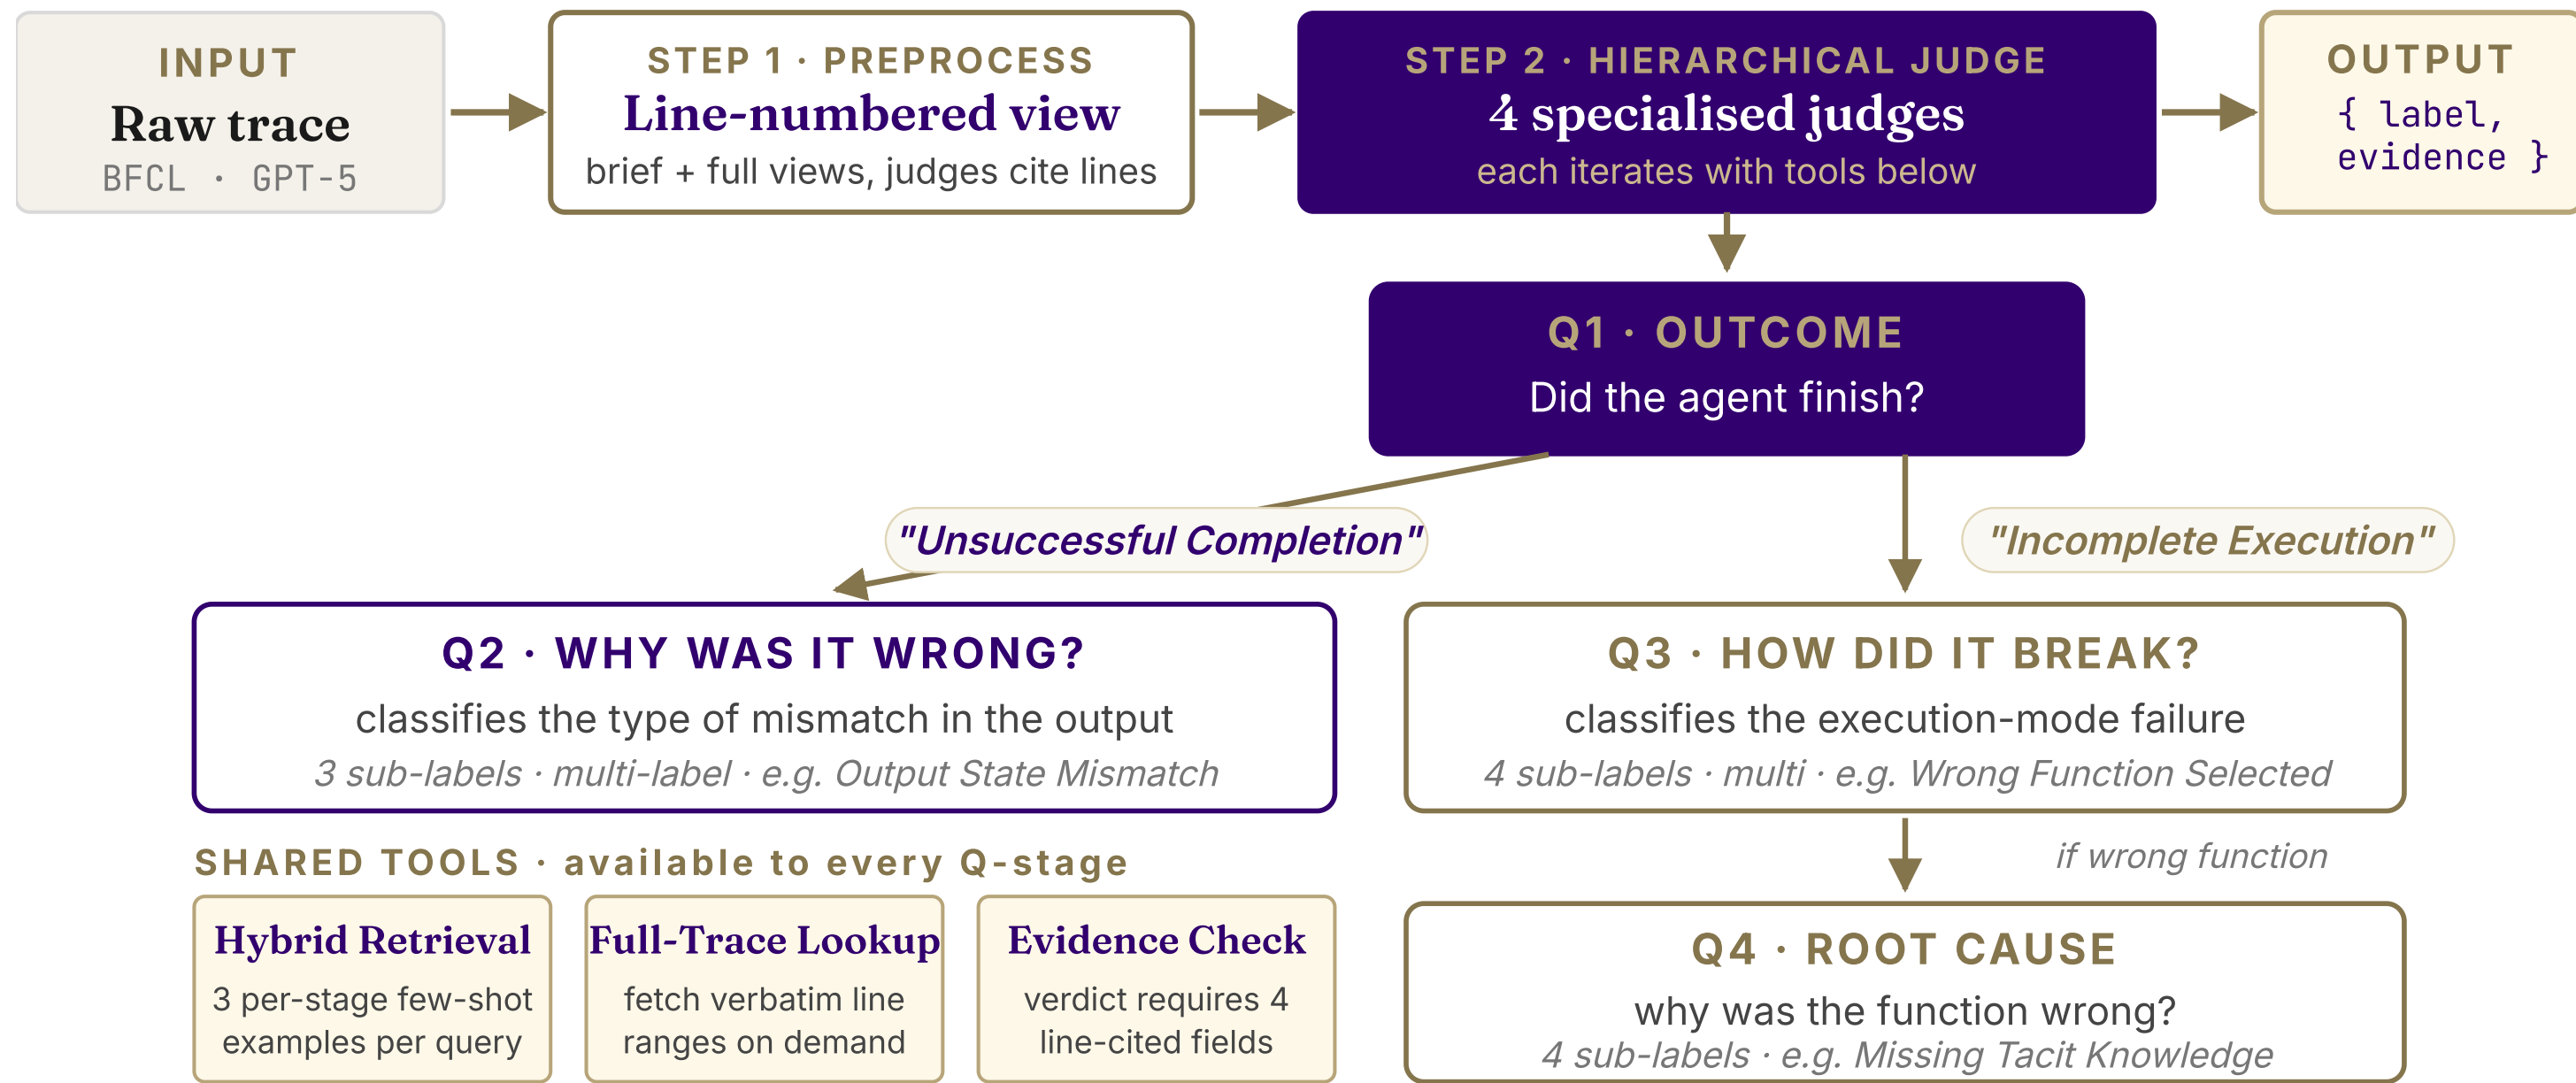


AFT is the answer key. Our judge is graded on how well its labels agree with the 500+ human annotations. Figure adapted from AFT taxonomy materials.

PIPELINE

Hierarchical Classification Pipeline

One model assigning 15 overlapping labels at once is hard. We split the decision into a chain of narrow questions, each with its own prompt, examples, and a shared toolset for fetching evidence.



RESULTS · TWO EXPERIMENTAL AXES

Architecture + Q3 prompt design.

Two independent levers: one changes the judge's *structure*, the other the *prompt* for a single stage. Each row adds one component to the row above.

AXIS 1 — ARCHITECTURE STACK (MICRO-F1)

CONFIGURATION	F1	Δ
Single-pass judge	0.557	—
+ Random few-shot	0.583	+ .026
+ Similarity few-shot	0.610	+ .027
+ Line-cited views	0.623	+ .013
+ Hierarchy (Q1-Q4)	0.706	+ .083

The four-way split does the heavy lifting: an 8-point jump, larger than every prompting trick combined. Few-shot barely helps until the structure is there.

AXIS 2 — Q3 PROMPT RUBRIC (Q3-STAGE MACRO-F1)

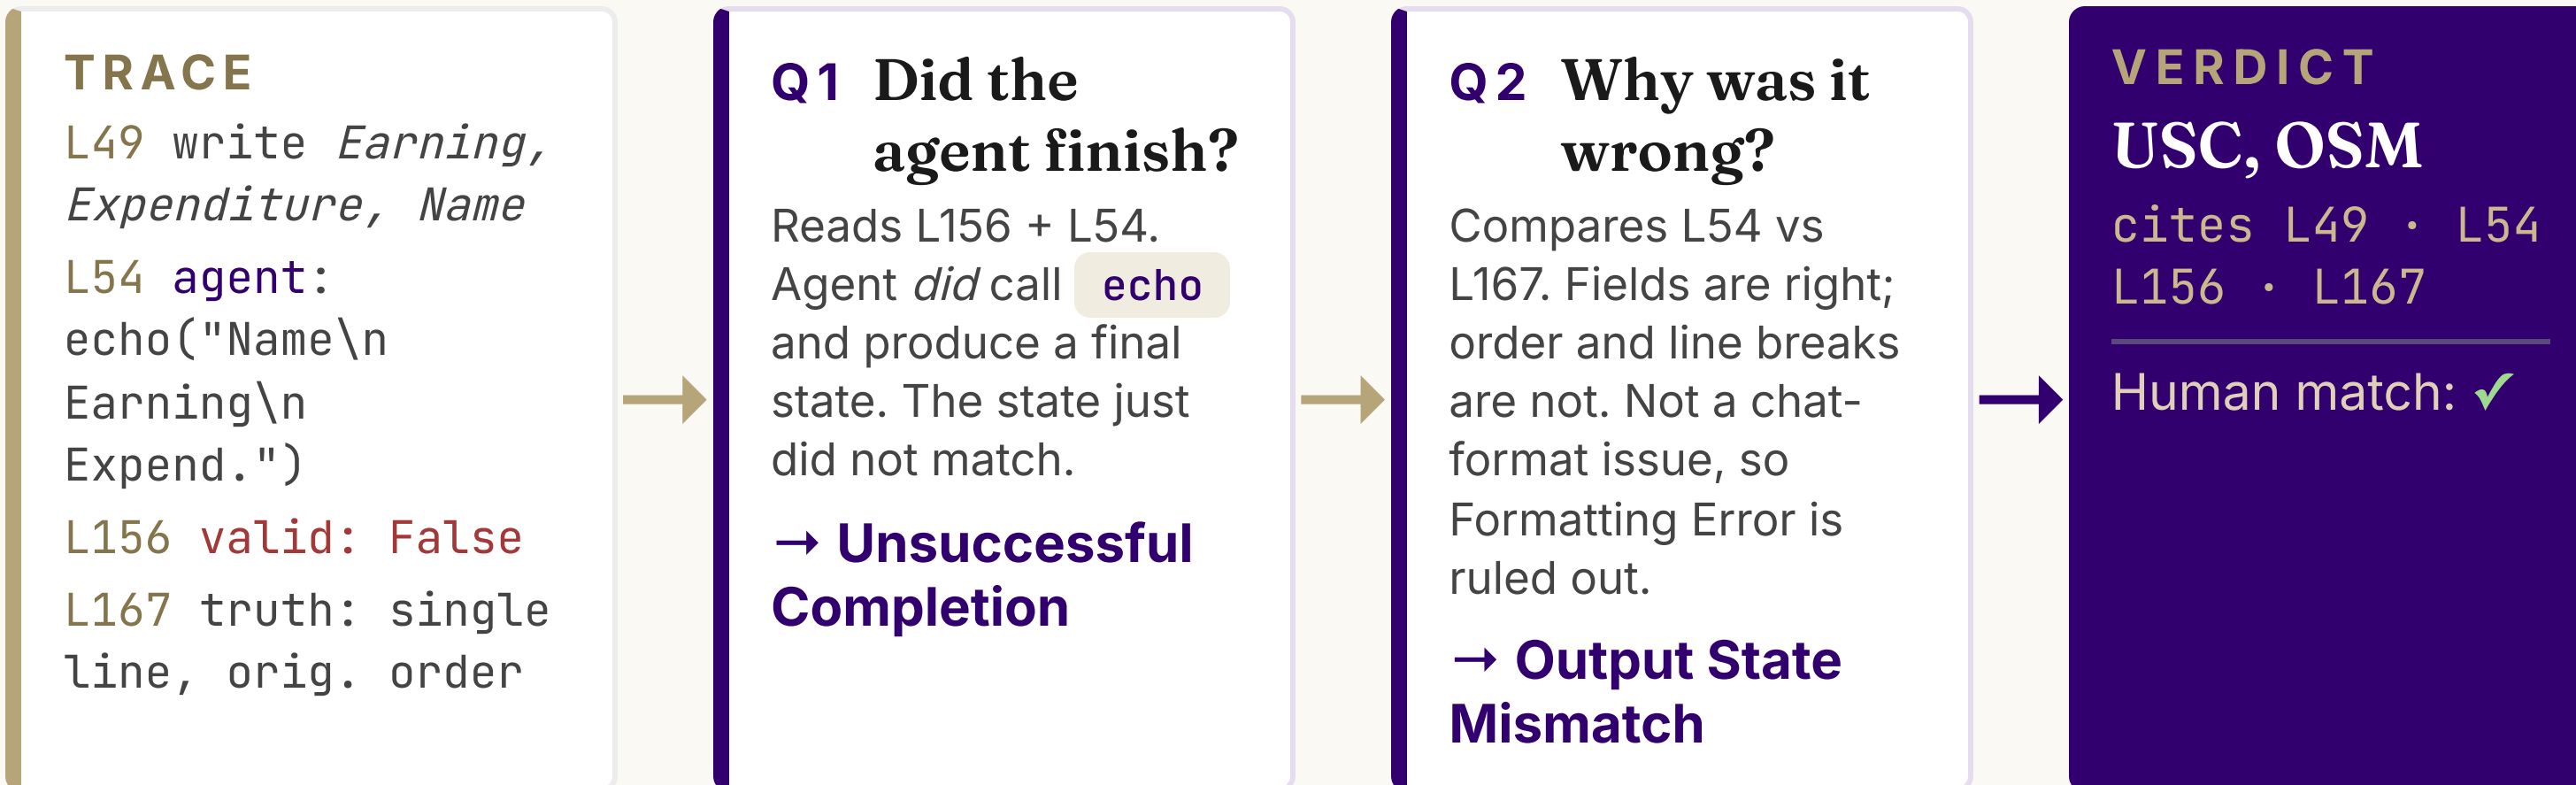
Q3 PROMPT	F1	Δ
Plain label list	0.603	—
+ Exclusion clauses	0.665	+ .062
+ Contrastive examples	0.722	+ .057
+ Chain-of-thought	0.771	+ .049

Telling the judge when *not* to apply a label, plus near-miss examples, nearly doubles the score on the four reasoning-error labels. Each change compounds.

EXAMPLE

How our pipeline labelled the *Annual Report* failure.

Trace **multi_turn_base_6** · BFCL · GPT-5. The judge walks Q1 → Q2 and submits a verdict citing specific lines.



DISCUSSION

Main Takeaways

- Structure matters more than prompting.** Just splitting the verdict into four staged questions raised micro-F1 from 0.56 to 0.71, before we tuned any prompt. The few-shot tricks added afterward were worth only 0.03 more.
- Once the structure is fixed, the rubric is the lever.** Telling the judge when *not* to apply a label, plus near-miss examples, took Q3 from 0.60 to 0.77 — a bigger jump than any architecture change.
- Some labels just won't budge.** We threw five different fixes at Skipped Sub-Task; every one landed at F1 ≈ 0.43. The judge keeps over-applying it, and prompt wording can't stop it.

PER-CLASS RESULTS

Where the judge wins and struggles.

F1 on the BFCL hold-out (35 traces, 142 labels).

TOP PERFORMERS

Agent Errors (parent)	1.00	35
Incomplete Execution	0.88	19
Unsuccessful Completion	0.83	16
Function Invoc. Error	0.80	3

WHERE IT STRUGGLES

Output State Mismatch	0.42	8
Skipped Sub-Task (plateau)	0.38	5
Function Docs Ignored (n=1)	0.00	1

NEXT STEPS

From a better judge to a scaled AFT study.

- Generalise beyond BFCL** — extend trace adapters to AppWorld and MCP-Universe.
- Cross-model evaluation** — re-run on Sonnet 4.5, Kimi K2, Qwen3-32B.
- Detect Fabrication structurally.** Prompt-only fixes have plateaued; a check on tool-call arguments is next.

CITATIONS

- [1] Cemri et al. (2025). *Why Do Multi-Agent LLM Systems Fail?* arXiv:2503.13657.
- [2] Zhu et al. (2025). *Where LLM Agents Fail and How They Can Learn From Failures.* arXiv:2509.25370.
- [3] Zhang & Agrawala (2026). *View-oriented Conversation Compiler for Agent Trace Analysis.* arXiv:2603.29678.